

telecom-customer-churn-prediction

October 9, 2023

1 Telecom Customer Churn Prediction

The aim of this project is to analyze customer demographics, services, tenure and other variables to predict whether a particular customer will churn or not.

2 Import libraries :-

Data Analysis and visualization libraries :

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
# sns.set_style("whitegrid")
import warnings
warnings.filterwarnings('ignore')
np.random.seed(42)
```

```
[2]: df=pd.read_csv("/content/WA_Fn-UseC_-Telco-Customer-Churn.xls")
df
```

```
[2]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	
	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\	

0	No	No phone service	DSL	No	...
1	Yes	No	DSL	Yes	...
2	Yes	No	DSL	Yes	...
3	No	No phone service	DSL	Yes	...
4	Yes	No	Fiber optic	No	...
...
7038	Yes	Yes	DSL	Yes	...
7039	Yes	Yes	Fiber optic	No	...
7040	No	No phone service	DSL	Yes	...
7041	Yes	Yes	Fiber optic	No	...
7042	Yes	No	Fiber optic	Yes	...

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	
2	No	No	No	No	Month-to-month	
3	Yes	Yes	No	No	One year	
4	No	No	No	No	Month-to-month	
...	
7038	Yes	Yes	Yes	Yes	One year	
7039	Yes	No	Yes	Yes	One year	
7040	No	No	No	No	Month-to-month	
7041	No	No	No	No	Month-to-month	
7042	Yes	Yes	Yes	Yes	Two year	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.5	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	
4	Yes	Electronic check	70.70	151.65	
...	
7038	Yes	Mailed check	84.80	1990.5	
7039	Yes	Credit card (automatic)	103.20	7362.9	
7040	Yes	Electronic check	29.60	346.45	
7041	Yes	Mailed check	74.40	306.6	
7042	Yes	Bank transfer (automatic)	105.65	6844.5	

	Churn
0	No
1	No
2	Yes
3	No
4	Yes
...	...
7038	No
7039	No

```
7040    No
7041   Yes
7042    No
```

```
[7043 rows x 21 columns]
```

```
[3]: df.shape
```

```
[3]: (7043, 21)
```

```
[4]: df.head()
```

```
[4]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female                0     Yes           No         1           No
1  5575-GNVDE   Male                0     No            No        34           Yes
2  3668-QPYBK   Male                0     No            No         2           Yes
3  7795-CFOCW   Male                0     No            No        45           No
4  9237-HQITU  Female                0     No            No         2           Yes
```

```
   MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service              DSL              No  ...              No
1                No              DSL              Yes  ...              Yes
2                No              DSL              Yes  ...              No
3  No phone service              DSL              Yes  ...              Yes
4                No  Fiber optic              No  ...              No
```

```
   TechSupport  StreamingTV  StreamingMovies  Contract  PaperlessBilling  \
0           No           No              No  Month-to-month              Yes
1           No           No              No    One year              No
2           No           No              No  Month-to-month              Yes
3           Yes           No              No    One year              No
4           No           No              No  Month-to-month              Yes
```

```
   PaymentMethod  MonthlyCharges  TotalCharges  Churn
0  Electronic check           29.85           29.85   No
1    Mailed check           56.95          1889.5   No
2    Mailed check           53.85           108.15  Yes
3  Bank transfer (automatic)       42.30          1840.75   No
4    Electronic check           70.70           151.65  Yes
```

```
[5 rows x 21 columns]
```

```
[5]: df.tail()
```

```
[5]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  \
7038  6840-RESVB   Male                0     Yes           Yes        24
7039  2234-XADUH  Female                0     Yes           Yes        72
```

7040	4801-JZAZL	Female	0	Yes	Yes	11
7041	8361-LTMKD	Male	1	Yes	No	4
7042	3186-AJIEK	Male	0	No	No	66

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
7038	Yes	Yes	DSL	Yes	...	
7039	Yes	Yes	Fiber optic	No	...	
7040	No	No phone service	DSL	Yes	...	
7041	Yes	Yes	Fiber optic	No	...	
7042	Yes	No	Fiber optic	Yes	...	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
7038	Yes	Yes	Yes	Yes	One year	
7039	Yes	No	Yes	Yes	One year	
7040	No	No	No	No	Month-to-month	
7041	No	No	No	No	Month-to-month	
7042	Yes	Yes	Yes	Yes	Two year	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
7038	Yes	Mailed check	84.80	1990.5	
7039	Yes	Credit card (automatic)	103.20	7362.9	
7040	Yes	Electronic check	29.60	346.45	
7041	Yes	Mailed check	74.40	306.6	
7042	Yes	Bank transfer (automatic)	105.65	6844.5	

	Churn
7038	No
7039	No
7040	No
7041	Yes
7042	No

[5 rows x 21 columns]

```
[6]: df.dtypes
```

```
[6]: customerID      object
gender              object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
OnlineBackup       object
```

```

DeviceProtection    object
TechSupport         object
StreamingTV         object
StreamingMovies     object
Contract            object
PaperlessBilling    object
PaymentMethod       object
MonthlyCharges      float64
TotalCharges        object
Churn               object
dtype: object

```

```
[7]: df.describe().T
```

```

[7]:
      count      mean      std   min  25%   50%   75%  \
SeniorCitizen  7043.0  0.162147  0.368612  0.00  0.0  0.00  0.00
tenure         7043.0  32.371149  24.559481  0.00  9.0  29.00  55.00
MonthlyCharges 7043.0  64.761692  30.090047  18.25 35.5  70.35  89.85

      max
SeniorCitizen    1.00
tenure           72.00
MonthlyCharges  118.75

```

Statistical description of Categorical features

```
[8]: df.describe(include='object')
```

```

[8]:
      customerID gender Partner Dependents PhoneService MultipleLines  \
count          7043    7043    7043        7043          7043        7043
unique          7043         2         2         2           2         3
top    7590-VHVEG    Male      No      No        Yes        No
freq           1    3555    3641    4933        6361        3390

      InternetService OnlineSecurity OnlineBackup DeviceProtection  \
count          7043          7043        7043          7043
unique           3           3         3         3
top    Fiber optic      No      No      No
freq          3096        3498        3088        3095

      TechSupport StreamingTV StreamingMovies      Contract  \
count          7043        7043        7043        7043
unique           3         3         3         3
top           No      No      No  Month-to-month
freq          3473        2810        2785        3875

      PaperlessBilling      PaymentMethod TotalCharges Churn

```

count	7043	7043	7043	7043
unique	2	4	6531	2
top	Yes	Electronic check		No
freq	4171	2365	11	5174

```
[9]: df.duplicated().sum()
```

```
[9]: 0
```

```
[10]: #covert 'TotalCharges' in to numerical datatype(some row contains " ")
df['TotalCharges']=df['TotalCharges'].replace(" ",np.nan)
df['TotalCharges']=pd.to_numeric(df['TotalCharges'])
```

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7032 non-null   float64
20  Churn                  7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

```
[12]: #checking for null values
```

```
df.isnull().sum()
```

```
[12]: customerID      0
      gender          0
      SeniorCitizen  0
      Partner        0
      Dependents     0
      tenure         0
      PhoneService   0
      MultipleLines  0
      InternetService 0
      OnlineSecurity 0
      OnlineBackup   0
      DeviceProtection 0
      TechSupport    0
      StreamingTV    0
      StreamingMovies 0
      Contract       0
      PaperlessBilling 0
      PaymentMethod  0
      MonthlyCharges 0
      TotalCharges   11
      Churn          0
      dtype: int64
```

11 missing datapoints can be observed in the Total Changes column

```
[13]: #Removal of rows with missing data
      df.dropna(inplace=True)
      df.reset_index()
```

```
[13]:
```

	index	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	0	7590-VHVEG	Female	0	Yes	No	1	
1	1	5575-GNVDE	Male	0	No	No	34	
2	2	3668-QPYBK	Male	0	No	No	2	
3	3	7795-CFOCW	Male	0	No	No	45	
4	4	9237-HQITU	Female	0	No	No	2	
...	
7027	7038	6840-RESVB	Male	0	Yes	Yes	24	
7028	7039	2234-XADUH	Female	0	Yes	Yes	72	
7029	7040	4801-JZAZL	Female	0	Yes	Yes	11	
7030	7041	8361-LTMKD	Male	1	Yes	No	4	
7031	7042	3186-AJIEK	Male	0	No	No	66	

	PhoneService	MultipleLines	InternetService	...	DeviceProtection	\
0	No	No phone service	DSL	...	No	
1	Yes	No	DSL	...	Yes	
2	Yes	No	DSL	...	No	
3	No	No phone service	DSL	...	Yes	

4	Yes	No	Fiber optic	...	No
...
7027	Yes	Yes	DSL	...	Yes
7028	Yes	Yes	Fiber optic	...	Yes
7029	No	No phone service	DSL	...	No
7030	Yes	Yes	Fiber optic	...	No
7031	Yes	No	Fiber optic	...	Yes

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	
...	
7027	Yes	Yes	Yes	One year	Yes	
7028	No	Yes	Yes	One year	Yes	
7029	No	No	No	Month-to-month	Yes	
7030	No	No	No	Month-to-month	Yes	
7031	Yes	Yes	Yes	Two year	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.50	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes
...
7027	Mailed check	84.80	1990.50	No
7028	Credit card (automatic)	103.20	7362.90	No
7029	Electronic check	29.60	346.45	No
7030	Mailed check	74.40	306.60	Yes
7031	Bank transfer (automatic)	105.65	6844.50	No

[7032 rows x 22 columns]

```
[14]: df.isnull().sum()
```

```
[14]: customerID      0
      gender          0
      SeniorCitizen  0
      Partner         0
      Dependents      0
      tenure          0
      PhoneService    0
      MultipleLines    0
      InternetService  0
```


OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0

dtype: int64

[15]: *#checking for duplicate values*

```
df.duplicated().sum()
```

[15]: 0

[16]: *#checking number of unique value in each column*

```
df.nunique()
```

[16]:

customerID	7032
gender	2
SeniorCitizen	2
Partner	2
Dependents	2
tenure	72
PhoneService	2
MultipleLines	3
InternetService	3
OnlineSecurity	3
OnlineBackup	3
DeviceProtection	3
TechSupport	3
StreamingTV	3
StreamingMovies	3
Contract	3
PaperlessBilling	2
PaymentMethod	4
MonthlyCharges	1584
TotalCharges	6530
Churn	2

dtype: int64

```
[17]: #to print unique value of each column
```

```
cols=df.columns
```

```
for i in cols:
```

```
    print(i, df[i].unique(), '\n')
```

```
customerID ['7590-VHVEG' '5575-GNVDE' '3668-QPYBK' ... '4801-JAZL' '8361-LTMKD'
'3186-AJIEK']
```

```
gender ['Female' 'Male']
```

```
SeniorCitizen [0 1]
```

```
Partner ['Yes' 'No']
```

```
Dependents ['No' 'Yes']
```

```
tenure [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27
  5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
 32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26 39]
```

```
PhoneService ['No' 'Yes']
```

```
MultipleLines ['No phone service' 'No' 'Yes']
```

```
InternetService ['DSL' 'Fiber optic' 'No']
```

```
OnlineSecurity ['No' 'Yes' 'No internet service']
```

```
OnlineBackup ['Yes' 'No' 'No internet service']
```

```
DeviceProtection ['No' 'Yes' 'No internet service']
```

```
TechSupport ['No' 'Yes' 'No internet service']
```

```
StreamingTV ['No' 'Yes' 'No internet service']
```

```
StreamingMovies ['No' 'Yes' 'No internet service']
```

```
Contract ['Month-to-month' 'One year' 'Two year']
```

```
PaperlessBilling ['Yes' 'No']
```

```
PaymentMethod ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
```

```
MonthlyCharges [29.85 56.95 53.85 ... 63.1  44.2  78.7 ]
```

```
TotalCharges [ 29.85 1889.5 108.15 ... 346.45 306.6 6844.5 ]
```

```
Churn ['No' 'Yes']
```

```
[18]: df['PhoneService'].value_counts()
```

```
[18]: Yes      6352  
      No       680  
      Name: PhoneService, dtype: int64
```

```
[19]: df['MultipleLines'].value_counts()
```

```
[19]: No              3385  
      Yes            2967  
      No phone service    680  
      Name: MultipleLines, dtype: int64
```

```
[20]: df['InternetService'].value_counts()
```

```
[20]: Fiber optic    3096  
      DSL          2416  
      No           1520  
      Name: InternetService, dtype: int64
```

```
[21]: df['OnlineSecurity'].value_counts()
```

```
[21]: No              3497  
      Yes            2015  
      No internet service  1520  
      Name: OnlineSecurity, dtype: int64
```

```
[22]: #create a list for each object column
```

```
lst=df.select_dtypes(include='object').columns.tolist()  
lst
```

```
[22]: ['customerID',  
      'gender',  
      'Partner',  
      'Dependents',  
      'PhoneService',  
      'MultipleLines',  
      'InternetService',  
      'OnlineSecurity',  
      'OnlineBackup',  
      'DeviceProtection',
```

```
'TechSupport',
'StreamingTV',
'StreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod',
'Churn']
```

```
[23]: #Uploading wrongly labeled data points : some data cells has 'No phone service'
      ↪and : 'No ,internet service' instead of 'No'
for i in lst:
    df[i]=df[i].replace('No phone service','No')
    df[i]=df[i].replace('No internet service','No')
```

```
[24]: df
```

```
[24]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
0	No	No	DSL	No	...	
1	Yes	No	DSL	Yes	...	
2	Yes	No	DSL	Yes	...	
3	No	No	DSL	Yes	...	
4	Yes	No	Fiber optic	No	...	
...	
7038	Yes	Yes	DSL	Yes	...	
7039	Yes	Yes	Fiber optic	No	...	
7040	No	No	DSL	Yes	...	
7041	Yes	Yes	Fiber optic	No	...	
7042	Yes	No	Fiber optic	Yes	...	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	
2	No	No	No	No	Month-to-month	
3	Yes	Yes	No	No	One year	

4	No	No	No	No	Month-to-month
...
7038	Yes	Yes	Yes	Yes	One year
7039	Yes	No	Yes	Yes	One year
7040	No	No	No	No	Month-to-month
7041	No	No	No	No	Month-to-month
7042	Yes	Yes	Yes	Yes	Two year

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.50	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	
4	Yes	Electronic check	70.70	151.65	
...	
7038	Yes	Mailed check	84.80	1990.50	
7039	Yes	Credit card (automatic)	103.20	7362.90	
7040	Yes	Electronic check	29.60	346.45	
7041	Yes	Mailed check	74.40	306.60	
7042	Yes	Bank transfer (automatic)	105.65	6844.50	

	Churn
0	No
1	No
2	Yes
3	No
4	Yes
...	...
7038	No
7039	No
7040	No
7041	Yes
7042	No

[7032 rows x 21 columns]

3 Exploratory Data Analysis


Visualization Of Data

for better understanding of data and to find relationship between dependent variables and the target variable

```
[25]: #distribution of customer tenure
      # plt.figure(figsize=(12,4))
```

```
sns.histplot(df['tenure'],color='teal',edgecolor='black',bins=20,kde=True)
plt.title('Customer Tenure Distribution')
plt.xlabel('Months')
plt.ylabel('Customer count')
```

[25]: Text(0, 0.5, 'Customer count')




[26]: df['tenure']

```
[26]: 0      1
      1     34
      2      2
      3     45
      4      2
      ..
     7038    24
     7039    72
     7040    11
     7041      4
     7042    66
```

Name: tenure, Length: 7032, dtype: int64


```
[27]: #distribution of monthly charges
plt.figure(figsize=(12,4))
sns.
    ↪histplot(df['MonthlyCharges'],color='teal',edgecolor='black',bins=20,kde=True)
plt.title('Monthly Charges Distribution')
plt.xlabel('Monthly charge')
plt.ylabel('Count')
```

[27]: Text(0, 0.5, 'Count')



```
[28]: #distribution of total charges
plt.figure(figsize=(12,4))
sns.histplot(df['TotalCharges'],color='teal',edgecolor='black',bins=20,kde=True)
plt.title('Total Charges Distribution')
plt.xlabel('Total charge')
plt.ylabel('Count')
```


[28]: Text(0, 0.5, 'Count')



```
[29]: #customer gender distribution
plt.figure(figsize=(10,6))
color=['teal','green']
label=['male','female']
plt.pie(df['gender'].value_counts(),labels=label,colors=color,autopct='%1.2f%%')
plt.title('Customer gender distribution')

plt.legend()
```

```
[29]: <matplotlib.legend.Legend at 0x7df11d1e7be0>
```

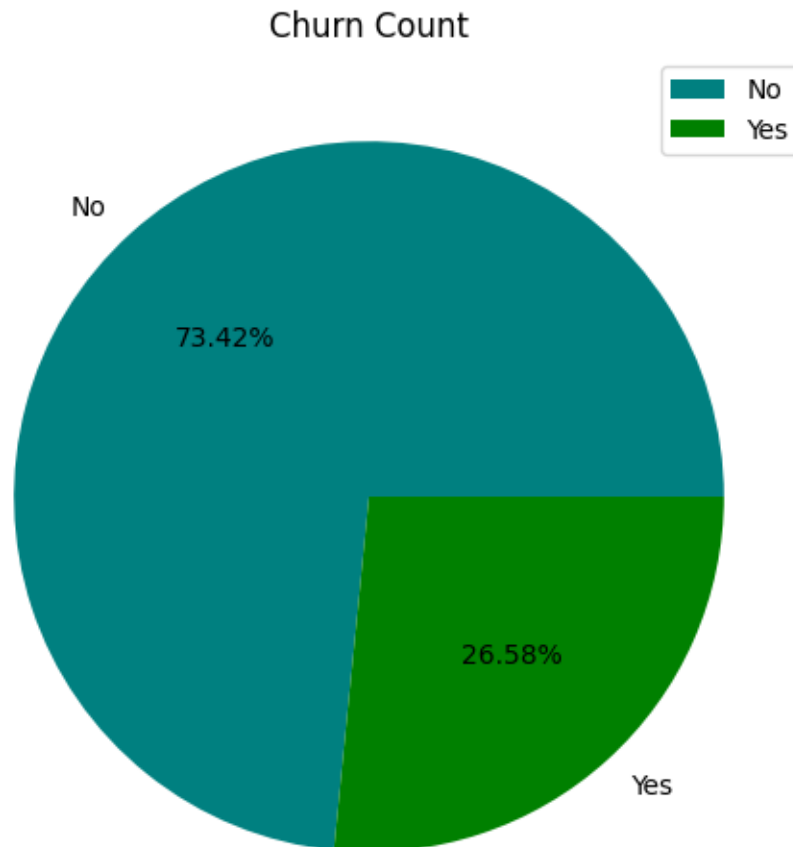



Gender distribution is approximately the same between males and females

COUNT OF CUSTOMER CHURN

```
[30]: plt.figure(figsize=(14,6))
color=['teal','green']
plt.pie(x= df['Churn'].value_counts(),labels= df['Churn'].
        ↪unique(),colors=color,autopct='%1.2f%%')
plt.title('Churn Count')
plt.legend()
```

```
[30]: <matplotlib.legend.Legend at 0x7df11d1a7070>
```



In the dataset, the number of churning customers is very less as compared to non churning. Only 26.49% churned from the telecom company. This could be a potential proof, that company is quite good at retaining its customers.

Customer Demographics and Churn

```
[31]: fig,ax=plt.subplots(2,2,figsize=(15,10))

#gender distribution
color=['teal','black']
sns.countplot(x='gender',data=df,hue='Churn',palette=color,ax=ax[0,0])
ax[0,0].set_title('Gender Distribution')

#senior citizen distriburion

sns.countplot(x='SeniorCitizen',data=df,hue='Churn',palette=color,ax=ax[0,1])
ax[0,1].set_title('SeniorCitizen And Churn')
```

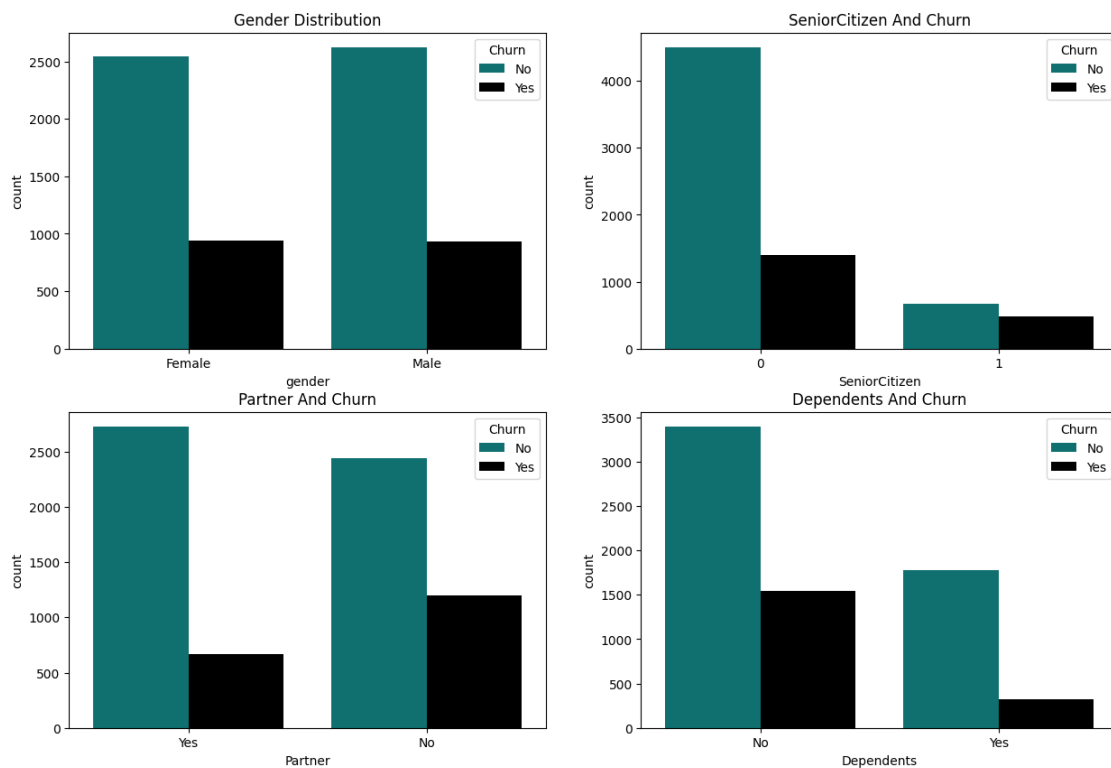
```
#partner distribution
```

```
sns.countplot(x='Partner',data=df,hue='Churn',palette=palette,ax=ax[1,0])
ax[1,0].set_title('Partner And Churn')
```

```
#Dependents distribution
```

```
sns.countplot(x='Dependents',data=df,hue='Churn',palette=palette,ax=ax[1,1])
ax[1,1].set_title('Dependents And Churn')
```

```
[31]: Text(0.5, 1.0, 'Dependents And Churn')
```



from above graph it is clear that male and female have same churn count. However, the senior citizens have a lesser churn count as compared to non senior citizens, The customers with no partners have higher churn count as compared to customers with partners. customers with no dependents have higher churn count as compared to customers with dependents.

CUSTOMER SERVICES AND CHURN

These graphs visualizes the relation between customer churn based on services opted by the customer

```
[32]: fig,ax=plt.subplots(3,3,figsize=(20,20))
color=['turquoise','black']

#phone service
sns.countplot(x='PhoneService',data=df,hue='Churn',palette=color,ax=ax[0,0])
ax[0,0].set_title('Phone Service')

#multiplelines
sns.countplot(x='MultipleLines',data=df,hue='Churn',palette=color,ax=ax[0,1])
ax[0,1].set_title('MultipleLines')
#internet service
sns.countplot(x='InternetService',data=df,hue='Churn',palette=color,ax=ax[0,2])
ax[0,2].set_title('Internet Service')

#online security service
sns.countplot(x='OnlineSecurity',data=df,hue='Churn',palette=color,ax=ax[1,0])
ax[1,0].set_title('Online Security')

#online backup
sns.countplot(x='OnlineBackup',data=df,hue='Churn',palette=color,ax=ax[1,1])
ax[1,1].set_title('OnlineBackup')

#Device protection
sns.countplot(x='DeviceProtection',data=df,hue='Churn',palette=color,ax=ax[1,2])
ax[2,0].set_title('Device protection')

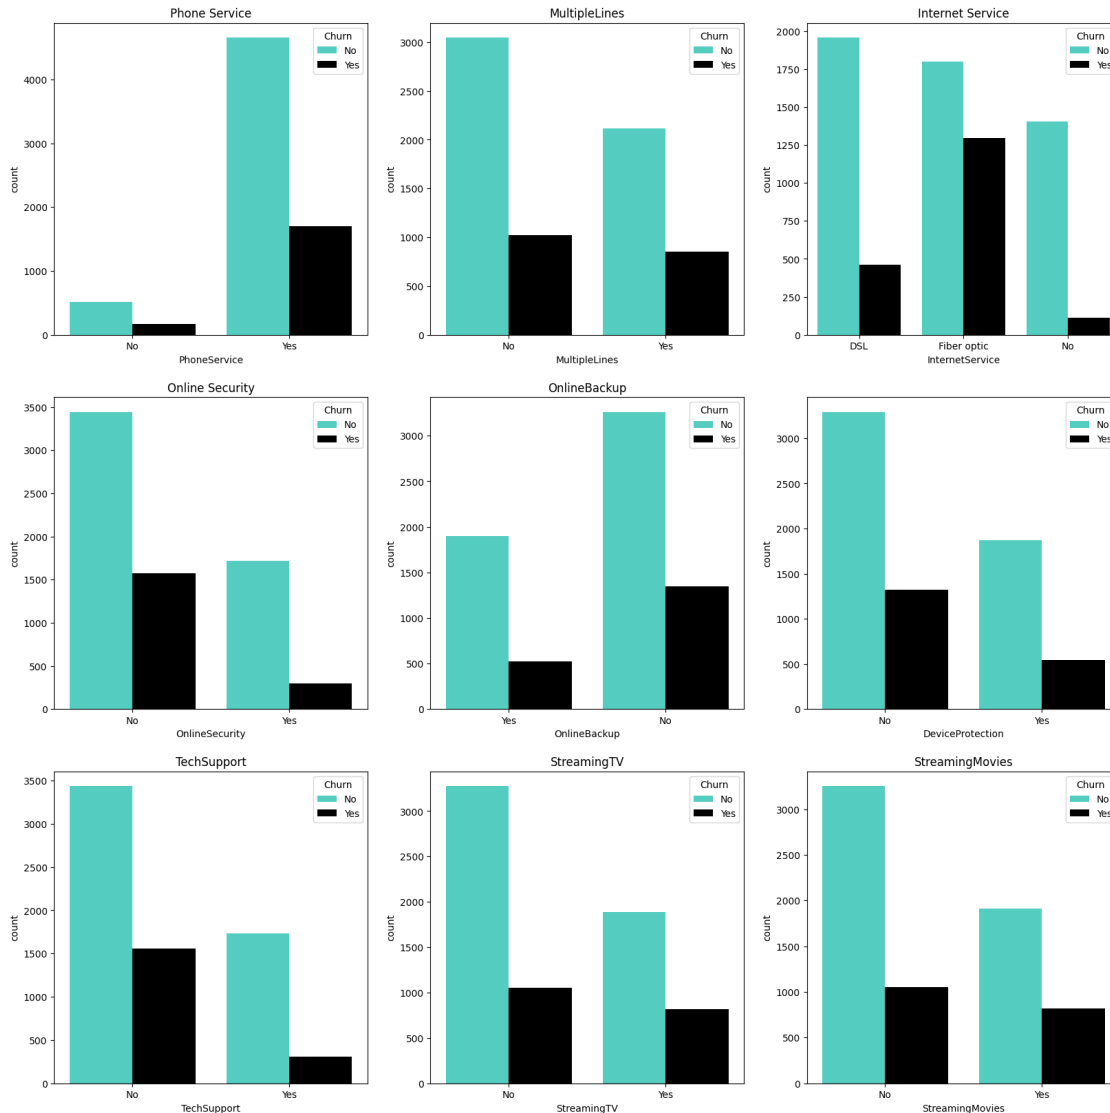
#tech support

sns.countplot(x='TechSupport',data=df,hue='Churn',palette=color,ax=ax[2,0])
ax[2,0].set_title('TechSupport')

#streaming TV
sns.countplot(x='StreamingTV',data=df,hue='Churn',palette=color,ax=ax[2,1])
ax[2,1].set_title('StreamingTV')

#Streaming Movies
sns.countplot(x='StreamingMovies',data=df,hue='Churn',palette=color,ax=ax[2,2])
ax[2,2].set_title('StreamingMovies')
```

```
[32]: Text(0.5, 1.0, 'StreamingMovies')
```




These graphs visualizes the relation between customer churn based on services opted by the customer. churn count is higher for the customers, who have taken multiple lines. the customers with streaming services have lower churn count

TENURE/CONTRACT AND CHURN

```
[33]: fig, ax = plt.subplots(1, 2, figsize=(15, 8))
sns.histplot(x = 'tenure', data = df, ax= ax[0], hue = 'Churn', multiple = 'stack').set_title('Cutomer Tenure (in Months) and Churn')
sns.countplot(x = 'Contract', data = df, ax= ax[1], hue = 'Churn').set_title('Contract Type and Churn')
```

```
[33]: Text(0.5, 1.0, 'Contract Type and Churn')
```




customer with long tenure have lowest churn rate, and customer with month to month contract have highest churn count than one or more year contract

Monthly Charges Distribution By Customer Churn

Charges And Churn

```
[35]: fig, ax = plt.subplots(1, 2, figsize=(15, 8))
fig.tight_layout(pad=5.0)
sns.histplot(x = 'MonthlyCharges', data = df, ax= ax[0], hue = 'Churn',
             multiple = 'stack').set_title('Monthly Charges')
sns.histplot(x = 'TotalCharges', data = df, ax= ax[1], hue = 'Churn', multiple =
             'stack').set_title('Total Charges')
```

```
[35]: Text(0.5, 1.0, 'Total Charges')
```



Customer with higher monthly charges have higher churn count, But customer with highest Total charges have the lower churn count. This could be possible when the customer has a long tenure and use lot of services. Therefore the company should focus on lowering the monthly charges in order to reduce churn count.

```
[36]: # Mean of monthly charges for churned and retained customers
mean_monthly_charges_churned = df[df['Churn'] == 'Yes']['MonthlyCharges'].mean()
mean_monthly_charges_retained = df[df['Churn'] == 'No']['MonthlyCharges'].mean()

print(f"Mean Monthly Charges for Churned Customers:␣
↪{mean_monthly_charges_churned:.2f}")
print(f"Mean Monthly Charges for Retained Customers:␣
↪{mean_monthly_charges_retained:.2f}")
```

Mean Monthly Charges for Churned Customers: 74.44

Mean Monthly Charges for Retained Customers: 61.31

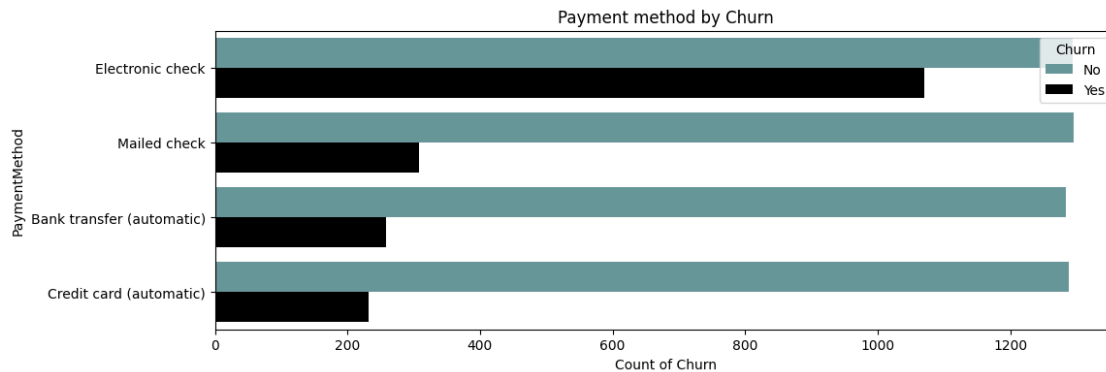
```
[37]: # Calculation of sum monthly charges by customer churn
monthly_charges_by_churn = df.groupby('Churn')['MonthlyCharges'].sum().
↪reset_index()
print(monthly_charges_by_churn)
```

	Churn	MonthlyCharges
0	No	316530.15
1	Yes	139130.85

Payment method by customer churn

```
[38]: # Payment method by customer churn
plt.figure(figsize=(12,4))
color=['cadetblue','black']
sns.countplot(data=df, y = 'PaymentMethod', hue = 'Churn',palette=color)
plt.xlabel('Count of Churn')
plt.title('Payment method by Churn')
```

```
[38]: Text(0.5, 1.0, 'Payment method by Churn')
```



Most of churn customers using electronic check for payment.

CORRELATION

Correlation in machine learning refers to the statistical relationship between two or more variables. It measures how closely the values of these variables are related to each other. Correlation is often used to understand the association between features (independent variables) and the target variable (dependent variable) in a dataset

```
[39]: df.corr()
```


```
[39]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
SeniorCitizen	1.000000	0.015683	0.219874	0.102411
tenure	0.015683	1.000000	0.246862	0.825880
MonthlyCharges	0.219874	0.246862	1.000000	0.651065
TotalCharges	0.102411	0.825880	0.651065	1.000000

GRAPHICAL REPRESENTATION OF CORRELATION(HEAT MAP)



```
[40]: sns.heatmap(df.corr(),annot=True)
```


```
[40]: <Axes: >
```

checking Outliers

```
[41]: lst1=['tenure', 'MonthlyCharges', 'TotalCharges']
      for i in lst1:
          sns.boxplot(df[i])
          plt.show()
```





```
[42]: df['customerID'].value_counts()
```

```
[42]: 7590-VHVEG      1
      0265-PSUAE      1
      2956-GGUCQ      1
      6008-NAIXK      1
      5956-YHHRX      1
      ..
      7874-ECPQJ      1
      9796-MVYXX      1
      2637-FKFSY      1
      1552-AAGRXX     1
      3186-AJIEK      1
      Name: customerID, Length: 7032, dtype: int64
```

```
[43]: #large no of unique values so drop customer ID
      df.drop(['customerID'],axis=1,inplace=True)
```

3.1 LABEL ENCODING

```
[44]: df['Contract'].value_counts()
```

```
[44]: Month-to-month      3875
      Two year           1685
      One year           1472
      Name: Contract, dtype: int64
```

```
[45]: from sklearn.preprocessing import LabelEncoder
      #columns for label encoding
      cols = df.columns[df.dtypes == 'object']
      #Label encoder object
      le = LabelEncoder()
      #Label encoding the columns
      for i in cols:
          le.fit(df[i])
          df[i] = le.transform(df[i])
          print(i, df[i].unique(), '\n')
```

```
gender [0 1]
```

```
Partner [1 0]
```

```
Dependents [0 1]
```

```
PhoneService [0 1]
```

```
MultipleLines [0 1]
```

```
InternetService [0 1 2]
```

```
OnlineSecurity [0 1]
```

```
OnlineBackup [1 0]
```

```
DeviceProtection [0 1]
```

```
TechSupport [0 1]
```

```
StreamingTV [0 1]
```

```
StreamingMovies [0 1]
```

```
Contract [0 1 2]
```

```
PaperlessBilling [1 0]
```

PaymentMethod [2 3 0 1]

Churn [0 1]

X AND Y SEPERATION

```
[46]: x=df.drop(columns='Churn',axis=1)
      y=df['Churn']
      x
```

```
[46]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	
...	
7038	1	0	1	1	24	1	
7039	0	0	1	1	72	1	
7040	0	0	1	1	11	0	
7041	1	1	1	0	4	1	
7042	1	0	0	0	66	1	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	0	0	0	1	
1	0	0	1	0	
2	0	0	1	1	
3	0	0	1	0	
4	0	1	0	0	
...	
7038	1	0	1	0	
7039	1	1	0	1	
7040	0	0	1	0	
7041	1	1	0	0	
7042	0	1	1	0	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	0	0	0	0	0	
1	1	0	0	0	1	
2	0	0	0	0	0	
3	1	1	0	0	1	
4	0	0	0	0	0	
...	
7038	1	1	1	1	1	
7039	1	0	1	1	1	
7040	0	0	0	0	0	
7041	0	0	0	0	0	

7042	1	1	1	1	2
	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	
0	1	2	29.85	29.85	
1	0	3	56.95	1889.50	
2	1	3	53.85	108.15	
3	0	0	42.30	1840.75	
4	1	2	70.70	151.65	
...	
7038	1	3	84.80	1990.50	
7039	1	1	103.20	7362.90	
7040	1	2	29.60	346.45	
7041	1	3	74.40	306.60	
7042	1	0	105.65	6844.50	

[7032 rows x 19 columns]

[47]: y

```
[47]: 0      0
      1      0
      2      1
      3      0
      4      1
      ..
      7038    0
      7039    0
      7040    0
      7041    1
      7042    0
      Name: Churn, Length: 7032, dtype: int64
```

FEATURE SCALING

```
[48]: from sklearn.preprocessing import MinMaxScaler
      mm=MinMaxScaler()
      sc_x=mm.fit_transform(x)
      sc_x
```

```
[48]: array([[0.          , 0.          , 1.          , ..., 0.66666667, 0.11542289,
              0.0012751 ],
              [1.          , 0.          , 0.          , ..., 1.          , 0.38507463,
              0.21586661],
              [1.          , 0.          , 0.          , ..., 1.          , 0.35422886,
              0.01031041],
              ...,
              [0.          , 0.          , 1.          , ..., 0.66666667, 0.11293532,
```

```

0.03780868],
[1.          , 1.          , 1.          , ..., 1.          , 0.55870647,
0.03321025],
[1.          , 0.          , 0.          , ..., 0.          , 0.86965174,
0.78764136]])

```

SPLIT DATA INTO TRAINING AND TESTING SETS

```

[49]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(sc_x,y,test_size=0.
↪3,random_state=42)

```

```

[50]: y.value_counts()

```

```

[50]: 0    5163
      1    1869
      Name: Churn, dtype: int64

```

Here, we can see that accuracy of the dataset is pretty good but this is an imbalanced data set. So the chance of false prediction in the class with lower value counts is high. Inorder to rectify that problem, we use the oversampling method

Balance dataset- Over Sampling

```

[51]: from imblearn.over_sampling import SMOTE
      smote = SMOTE(random_state=42)
      x_train_resampled, y_train_resampled = smote.fit_resample(x_train, y_train)

```

```

[52]: y_train_resampled.value_counts()

```

```

[52]: 1    3614
      0    3614
      Name: Churn, dtype: int64

```

```

[53]: mm=MinMaxScaler()
      sc_x1=mm.fit_transform(x_train_resampled)
      sc_x1

```

```

[53]: array([[0.          , 0.          , 0.          , ..., 1.          , 0.40109616,
0.06012036],
[1.          , 0.          , 0.          , ..., 0.66666667, 0.43148979,
0.18037261],
[0.          , 0.          , 0.          , ..., 0.66666667, 0.51519681,
0.02326346],
...,
[1.          , 0.          , 0.          , ..., 0.43142225, 0.88224269,
0.49283617],

```

```
[0.          , 0.          , 0.          , ..., 0.66666667, 0.60096037,
 0.05156806],
[1.          , 0.          , 0.05846709, ..., 1.          , 0.33049549,
 0.09687183]])
```

```
[54]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(sc_x1,y_train_resampled)
```

```
[ ]: # from sklearn.model_selection import train_test_split
#
↳x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=42,test_size=0.
↳20)
```

```
[ ]: # from sklearn.preprocessing import StandardScaler
# scaler=StandardScaler()
# scaler.fit(x_train)
# x_train=scaler.fit_transform(x_train)
# x_test=scaler.fit_transform(x_test)
```

3.2 MODEL BUILDING

K Nearest Neighbor Classifier

NAIVE BAYES

SVM

RANDOM FOREST CLASSIFIER

MODEL EVALUATION TOOLS

ACCURACY SCORE

CLASSIFICATION REPORT

CONFUSION MATRIX

importing classification report and accuracy score for model evaluation

```
[55]: from sklearn.metrics import
↳accuracy_score,classification_report,ConfusionMatrixDisplay
```

1.KNN

```
[56]: from sklearn.neighbors import KNeighborsClassifier

knn=KNeighborsClassifier()
```

Hyperparameter Tuning using GridSearchCV

```
[96]: from sklearn.model_selection import GridSearchCV
from sklearn.metrics.pairwise import distance
```



```

np.random.seed(42)

param={'n_neighbors':[1,3,5,7,9],
       'weights':['uniform','distance'],
       'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute']}
}
gsv_knn=GridSearchCV(knn,param,cv=10,scoring='accuracy')

gsv_knn.fit(x_train,y_train)

#Best parameters
print(gsv_knn.best_params_)

```

```
{'algorithm': 'ball_tree', 'n_neighbors': 1, 'weights': 'uniform'}
```

MODEL CREATION

```

[97]: knn=KNeighborsClassifier(algorithm='ball_tree',n_neighbors=1,weights='uniform')
      knn.fit(x_train,y_train)
      knn_pred=knn.predict(x_test)

```

MODEL EVALUATION

```

[59]: #accuracy score
      knn_acc=accuracy_score(y_test,knn_pred)
      knn_acc

```

```
[59]: 0.8206972883231876
```


```

[60]: #classification report and confusion matrix
      print(classification_report(y_test,knn_pred))
      print(ConfusionMatrixDisplay.from_predictions(y_test,knn_pred))

```

	precision	recall	f1-score	support
0	0.88	0.75	0.81	916
1	0.78	0.89	0.83	891
accuracy			0.82	1807
macro avg	0.83	0.82	0.82	1807
weighted avg	0.83	0.82	0.82	1807

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7df117fd2560>
```



2. SVM (SUPPORT VECTOR MACHINE)

```
[61]: from sklearn.svm import SVC
      svm=SVC()
```

Hyperparameter Tuning using GridSearchCV

```
[63]: param_grid_svm = {'C': [0.1, 1, 10],
      'kernel': ['linear', 'rbf'],
      'gamma': ['scale', 'auto', 0.1, 1]}

      grid_search_svm = GridSearchCV(svm, param_grid_svm, scoring='recall', cv=10)
      grid_search_svm.fit(x_train, y_train)

      print(grid_search_svm.best_params_)
```

```
{'C': 10, 'gamma': 1, 'kernel': 'rbf'}
```

MODEL CREATION

```
[67]: svm=SVC(C=10,gamma=1,kernel='rbf')
      svm.fit(x_train,y_train)
      svm_pred=svm.predict(x_test)
```

MODEL EVALUATION


```
[68]: #accuracy score
svm_acc=accuracy_score(y_test,svm_pred)
svm_acc
```

[68]: 0.8494742667404538

```
[69]: #classification report and confusion matrix
print(classification_report(y_test,svm_pred))
print(ConfusionMatrixDisplay.from_predictions(y_test,svm_pred))
```

	precision	recall	f1-score	support
0	0.89	0.81	0.84	916
1	0.82	0.89	0.85	891
accuracy			0.85	1807
macro avg	0.85	0.85	0.85	1807
weighted avg	0.85	0.85	0.85	1807

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7df115f335e0>



3. DECISION TREE CLASSIFIER

```
[70]: from sklearn.tree import DecisionTreeClassifier
      #Decision Tree Classifier Object
      dtree = DecisionTreeClassifier()
```

Hyperparameter Tuning using GridSearchCV

```
[72]: from sklearn.model_selection import GridSearchCV

      #parameter grid
      param_grid = {
          'max_depth': [2,4,6,8,10],
          'min_samples_leaf': [2,4,6,8,10],
          'min_samples_split': [2,4,6,8,10],
          'criterion': ['gini', 'entropy'],
          'random_state': [0,42]
      }

      #Grid Search Object with Decision Tree Classifier
      grid_search_d = GridSearchCV(estimator = dtree, param_grid = param_grid, cv = 3,
      ↪n_jobs = -1, verbose = 2, scoring='accuracy')

      #Fitting the data
      grid_search_d.fit(x_train, y_train)
```

Fitting 3 folds for each of 500 candidates, totalling 1500 fits

```
[72]: GridSearchCV(cv=3, estimator=DecisionTreeClassifier(), n_jobs=-1,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': [2, 4, 6, 8, 10],
                              'min_samples_leaf': [2, 4, 6, 8, 10],
                              'min_samples_split': [2, 4, 6, 8, 10],
                              'random_state': [0, 42]},
                  scoring='accuracy', verbose=2)
```

```
[73]: #Best parameters
      print(grid_search_d.best_params_)
```

```
{'criterion': 'entropy', 'max_depth': 8, 'min_samples_leaf': 8,
 'min_samples_split': 2, 'random_state': 42}
```

MODEL CREATION

```
[74]: dtree = DecisionTreeClassifier(criterion='entropy', max_depth=8,
      ↪min_samples_leaf=8,min_samples_split=2, random_state=42)
      dtree.fit(x_train,y_train)
```

```
dtree_pred=dtree.predict(x_test)
```

MODEL EVALUATION


```
[75]: dtree_acc=accuracy_score(y_test,dtree_pred)
dtree_acc
```

```
[75]: 0.8007747648035418
```

```
[76]: #classification report and confusion matrix
print(classification_report(y_test,dtree_pred))
print(ConfusionMatrixDisplay.from_predictions(y_test,dtree_pred))
```

	precision	recall	f1-score	support
0	0.81	0.79	0.80	916
1	0.79	0.81	0.80	891
accuracy			0.80	1807
macro avg	0.80	0.80	0.80	1807
weighted avg	0.80	0.80	0.80	1807

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7df115eef310>



4.RANDOM FOREST

```
[77]: from sklearn.ensemble import RandomForestClassifier
      rf=RandomForestClassifier()
      rf.get_params()
```

```
[77]: {'bootstrap': True,
      'ccp_alpha': 0.0,
      'class_weight': None,
      'criterion': 'gini',
      'max_depth': None,
      'max_features': 'sqrt',
      'max_leaf_nodes': None,
      'max_samples': None,
      'min_impurity_decrease': 0.0,
      'min_samples_leaf': 1,
      'min_samples_split': 2,
      'min_weight_fraction_leaf': 0.0,
      'n_estimators': 100,
      'n_jobs': None,
      'oob_score': False,
      'random_state': None,
      'verbose': 0,
      'warm_start': False}
```

Hyperparameter Tuning using GridSearchCV

```
[78]: param_rf = {'n_estimators': [50, 100, 150],
                  'max_depth': [ 10, 20, 30,40,50]
                }

      grid_search_rf = GridSearchCV(rf, param_rf, scoring = 'accuracy', cv=10)
      grid_search_rf.fit(x_train, y_train)

      print(grid_search_rf.best_params_)
```

```
{'max_depth': 50, 'n_estimators': 150}
```

MODEL CREATION

```
[83]: rf=RandomForestClassifier(max_depth=50,n_estimators=150)
      rf.fit(x_train,y_train)
      rf_pred=rf.predict(x_test)
```

MODEL EVALUATION


```
[84]: #accuracy score
rf_acc=accuracy_score(y_test,rf_pred)
rf_acc
```

```
[84]: 0.8472606530160487
```

```
[85]: #classification report and confusion matrix
print(classification_report(y_test,rf_pred))
print(ConfusionMatrixDisplay.from_predictions(y_test,rf_pred))
```

	precision	recall	f1-score	support
0	0.85	0.84	0.85	916
1	0.84	0.85	0.85	891
accuracy			0.85	1807
macro avg	0.85	0.85	0.85	1807
weighted avg	0.85	0.85	0.85	1807

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7df115eed7b0>



5. NAIVE BAYES CLASSIFIER

```
[86]: from sklearn.naive_bayes import GaussianNB
      nb=GaussianNB()
      nb.fit(x_train,y_train)
      nb_pred=nb.predict(x_test)
```

MODEL EVALUATION


```
[87]: nb_acc=(accuracy_score(y_test,nb_pred))
      nb_acc
```

```
[87]: 0.7714443829551744
```

```
[88]: #classification report and confusion matrix
      print(classification_report(y_test,nb_pred))
      print(ConfusionMatrixDisplay.from_predictions(y_test,nb_pred))
```

	precision	recall	f1-score	support
0	0.80	0.73	0.76	916
1	0.74	0.82	0.78	891
accuracy			0.77	1807
macro avg	0.77	0.77	0.77	1807
weighted avg	0.77	0.77	0.77	1807

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at
0x7df117fa5600>
```

MODEL EVALUATION - Accuracy_Score visualization


ACCURACY SCORE COMPARISON BETWEEN MODELS

```
[89]: model=['KNN', 'SVM', 'DECISION TREE', 'RANDOM FOREST', 'NAIVE BAYES']
accuracy_scores=[knn_acc,svm_acc,dtree_acc,rf_acc,nb_acc]
accuracy_scores
```

```
[89]: [0.8206972883231876,
0.8494742667404538,
0.8007747648035418,
0.8472606530160487,
0.7714443829551744]
```

```
[90]: color=['black', 'teal', 'turquoise', 'cadetblue', 'green']
plt.figure(figsize=(10,6))
sns.barplot(x=model,y=accuracy_scores,palette=color)
plt.ylabel('accuracy_scores')
plt.title('Accuracy scores of models')
# Adding percentage labels
for i, score in enumerate(accuracy_scores):
    plt.text(i, score +0.01, f'{score*100:.2f}%', ha = 'center')
```

```
plt.show()
```



By plotting accuracy score with percentage label, it is clear that support vector machine classifier and Random forest classifier shows highest accuracy score than other models.

3.3 Confusion Matrix Heatmap

A confusion matrix is a common tool used in machine learning to evaluate the performance of a classification algorithm.

```
[95]: # #from sklearn.metrics import confusion_matrix
# fig, ax = plt.subplots(2, 2, figsize=(10, 7))

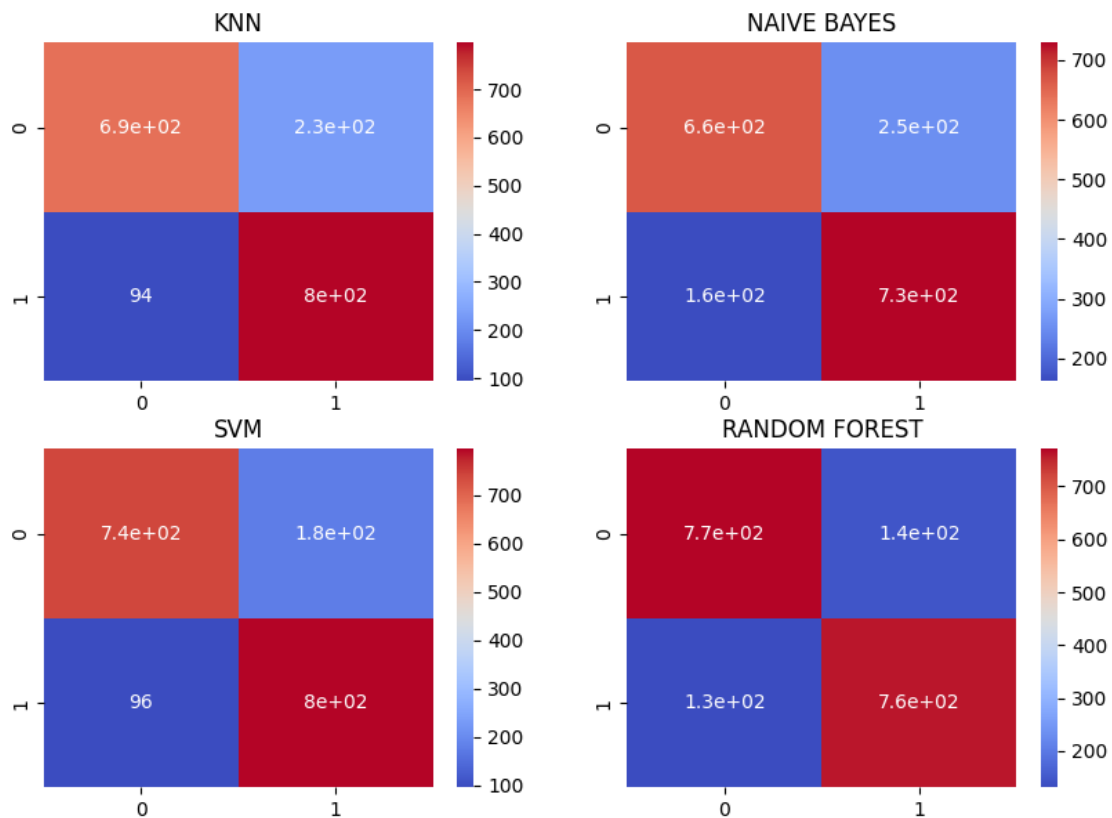
# #KNN Confusion Matrix
# sns.heatmap(confusion_matrix(y_test, knn_pred), annot=True, ax=ax[0,0],
#             cmap='coolwarm')
# ax[0,0].set_title('KNN')

# #NAIVE BAYES confusion matrix
# sns.heatmap(confusion_matrix(y_test, nb_pred), annot=True, ax=ax[0,1],
#             cmap='coolwarm')
# ax[0,1].set_title('NAIVE BAYES')

# #SVM Confusion Matrix
```

```
# sns.heatmap(confusion_matrix(y_test, svm_pred), annot=True, ax=ax[1,0],  
#             cmap='coolwarm')  
# ax[1,0].set_title('SVM')  
  
# #RANDOM FORESTConfusion Matrix  
# sns.heatmap(confusion_matrix(y_test, rf_pred), annot=True, ax=ax[1,1],  
#             cmap='coolwarm')  
# ax[1,1].set_title('RANDOM FOREST')
```

[95]: Text(0.5, 1.0, 'RANDOM FOREST')



The confusion matrix heatmaps visualize the true positive and true negative results from the machine learning model.

DISTRIBUTION PLOT(Y_TEST V/S Y_PREDICTED VALUES)

[93]: fig, ax = plt.subplots(2, 2, figsize=(12, 10))

```
#KNN
```

```

sns.distplot(y_test, hist=False, color="r", label="Actual Value", ax=ax[0,0]).
    ↪set_title('KNN')
sns.distplot(knn_pred, hist=False, color="b", label="Fitted Values" ,
    ↪ax=ax[0,0])

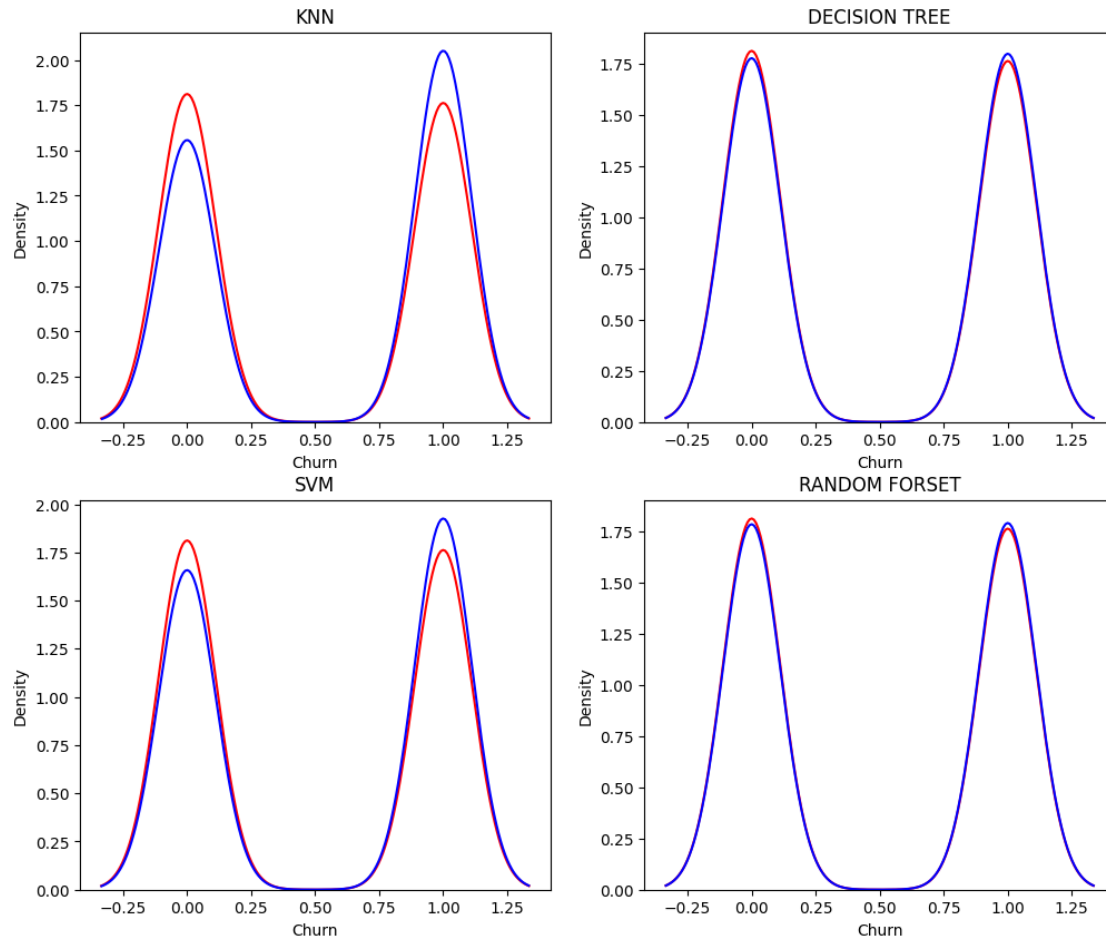
#DECISION TREE
sns.distplot(y_test, hist=False, color="r", label="Actual Value", ax=ax[0,1]).
    ↪set_title('DECISION TREE')
sns.distplot(dtrees_pred, hist=False, color="b", label="Fitted Values" ,
    ↪ax=ax[0,1])

#SVM
sns.distplot(y_test, hist=False, color="r", label="Actual Value", ax=ax[1,0]).
    ↪set_title('SVM')
sns.distplot(svm_pred, hist=False, color="b", label="Fitted Values" ,
    ↪ax=ax[1,0])

# #RANDOM FOREST
sns.distplot(y_test, hist=False, color="r", label="Actual Value", ax=ax[1,1]).
    ↪set_title('RANDOM FORSET')
sns.distplot(rf_pred, hist=False, color="b", label="Fitted Values" , ax=ax[1,1])

```

[93]: <Axes: title={'center': 'RANDOM FORSET'}, xlabel='Churn', ylabel='Density'>



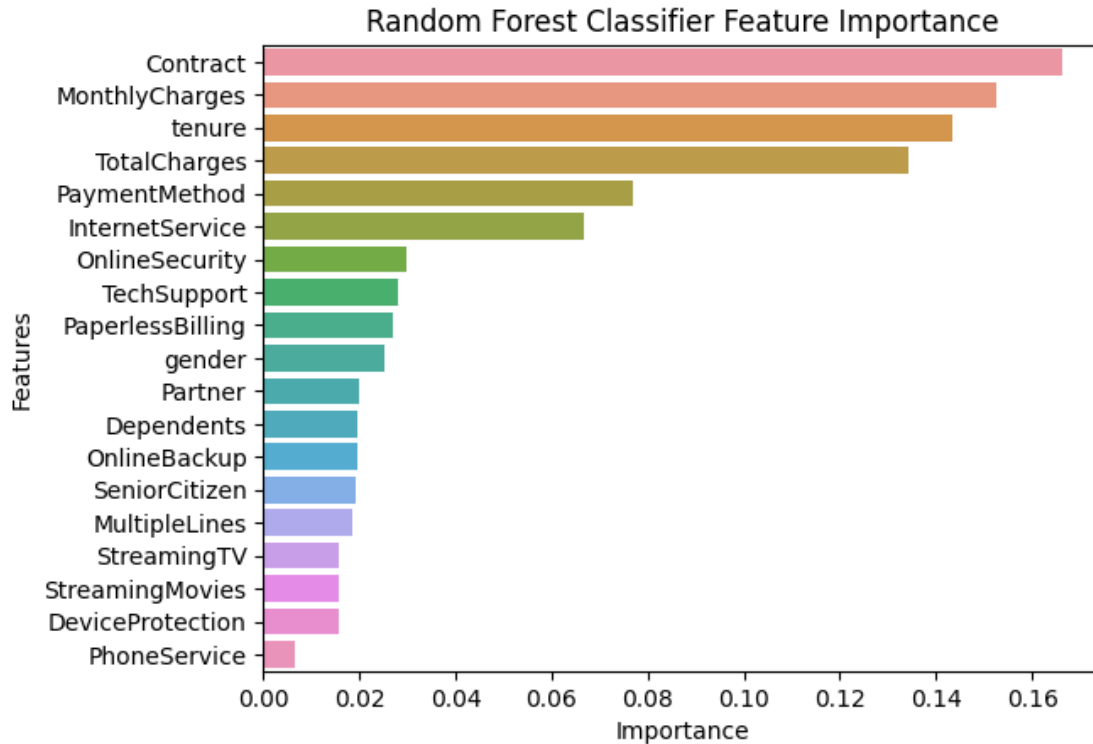
Distribution plots allow you to visually compare the distribution of your model's predictions (fitted values) with the distribution of the actual values. It provides a visual way to assess the goodness of fit of your model, understand the quality of predictions, and identify potential issues or areas for improvement.

Feature Importance

```
[94]: # Random Forest Classifier Feature Importance

x_train = pd.DataFrame(x_train, columns = x.columns)
feature_df = pd.DataFrame({'Features': x_train.columns, 'Importance': rf.
    ↪feature_importances_})
feature_df.sort_values('Importance', ascending=False, inplace=True)
sns.barplot(x = 'Importance', y = 'Features', data = feature_df).
    ↪set_title('Random Forest Classifier Feature Importance')
```

```
[94]: Text(0.5, 1.0, 'Random Forest Classifier Feature Importance')
```



Total charges, Monthly charges, tenure, contract are the most important features for predicting the customer churn hence company should focus on these features to reduce churn rate.

CONCLUSION

I have used five models - K Nearest Neighbor classifier, Naive Bayes classifier, Support Vector Machine (SVM), Decision Tree Classifier, Random Forest Classifier. The Support Vector Machine Classifier and Random Forest Classifier show the highest accuracy score and F1 Score. Therefore, the Support Vector Machine Classifier and Random Forest Classifier are good fits for predicting customer churn.