

# titanic-survival-prediction

January 20, 2024

## 1 TITANIC SURVIVAL PREDICTION

**AIM:** Building a model that predicts whether a passenger on the Titanic survived or not using the given dataset

### IMPORTING ESSENTIAL LIBRARIES

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
np.random.seed(42)
```

### DATA COLLECTION, PROCESSING AND ANALYSIS

```
[ ]: #takes a path to a CSV file and reads the data into a DataFrame
df=pd.read_csv("/content/tested.xls")
df
```

```
[ ]: PassengerId  Survived  Pclass  \
0             892         0       3
1             893         1       3
2             894         0       2
3             895         0       3
4             896         1       3
..          ...         ...     ...
413          1305         0       3
414          1306         1       1
415          1307         0       3
416          1308         0       3
417          1309         0       3
```

```
0              Name    Sex  Age  SibSp  Parch  \
0  Kelly, Mr. James  male  34.5     0     0
```

1	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0
2	Myles, Mr. Thomas Francis	male	62.0	0	0
3	Wirz, Mr. Albert	male	27.0	0	0
4	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1
..	...	...	...	...	...
413	Spector, Mr. Woolf	male	NaN	0	0
414	Oliva y Ocana, Dona. Fermina	female	39.0	0	0
415	Saether, Mr. Simon Sivertsen	male	38.5	0	0
416	Ware, Mr. Frederick	male	NaN	0	0
417	Peter, Master. Michael J	male	NaN	1	1

	Ticket	Fare	Cabin	Embarked
0	330911	7.8292	NaN	Q
1	363272	7.0000	NaN	S
2	240276	9.6875	NaN	Q
3	315154	8.6625	NaN	S
4	3101298	12.2875	NaN	S
..	...	...	...	...
413	A.5. 3236	8.0500	NaN	S
414	PC 17758	108.9000	C105	C
415	SOTON/O.Q. 3101262	7.2500	NaN	S
416	359309	8.0500	NaN	S
417	2668	22.3583	NaN	C

[418 rows x 12 columns]

```
[ ]: df.head()
```

```
[ ]: PassengerId  Survived  Pclass  \
0            892         0       3
1            893         1       3
2            894         0       2
3            895         0       3
4            896         1       3
```

	Name	Sex	Age	SibSp	Parch	\
0	Kelly, Mr. James	male	34.5	0	0	
1	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	
2	Myles, Mr. Thomas Francis	male	62.0	0	0	
3	Wirz, Mr. Albert	male	27.0	0	0	
4	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	

	Ticket	Fare	Cabin	Embarked
0	330911	7.8292	NaN	Q
1	363272	7.0000	NaN	S
2	240276	9.6875	NaN	Q
3	315154	8.6625	NaN	S

```
4 3101298 12.2875 NaN S
```

```
[ ]: df.tail()
```

```
[ ]:      PassengerId  Survived  Pclass                Name  Sex \
413          1305         0         3      Spector, Mr. Woolf  male
414          1306         1         1  Oliva y Ocana, Dona. Fermina  female
415          1307         0         3  Saether, Mr. Simon Sivertsen  male
416          1308         0         3      Ware, Mr. Frederick  male
417          1309         0         3  Peter, Master. Michael J  male
```

```
      Age  SibSp  Parch            Ticket     Fare Cabin Embarked
413   NaN     0     0          A.5. 3236     8.0500   NaN        S
414  39.0     0     0          PC 17758    108.9000  C105        C
415  38.5     0     0  SOTON/O.Q. 3101262     7.2500   NaN        S
416   NaN     0     0          359309     8.0500   NaN        S
417   NaN     1     1             2668     22.3583   NaN        C
```

```
[ ]: df.shape
```

```
[ ]: (418, 12)
```

```
[ ]: df.dtypes
```

```
[ ]: PassengerId    int64
Survived           int64
Pclass             int64
Name               object
Sex                object
Age                float64
SibSp              int64
Parch              int64
Ticket             object
Fare               float64
Cabin              object
Embarked           object
dtype: object
```

```
[ ]: df.describe().T
```

```
[ ]:      count      mean      std      min      25%      50% \
PassengerId  418.0  1100.500000  120.810458  892.00  996.2500  1100.5000
Survived      418.0    0.363636   0.481622    0.00    0.0000    0.0000
Pclass        418.0    2.265550   0.841838    1.00    1.0000    3.0000
Age           332.0   30.272590   14.181209    0.17   21.0000   27.0000
SibSp          418.0    0.447368   0.896760    0.00    0.0000    0.0000
Parch          418.0    0.392344   0.981429    0.00    0.0000    0.0000
```

Fare	417.0	35.627188	55.907576	0.00	7.8958	14.4542
------	-------	-----------	-----------	------	--------	---------

	75%	max
PassengerId	1204.75	1309.0000
Survived	1.00	1.0000
Pclass	3.00	3.0000
Age	39.00	76.0000
SibSp	1.00	8.0000
Parch	0.00	9.0000
Fare	31.50	512.3292

```
[ ]: df.describe(include='object').T
```

	count	unique	top	freq
Name	418	418	Kelly, Mr. James	1
Sex	418	2	male	266
Ticket	418	363	PC 17608	5
Cabin	91	76	B57 B59 B63 B66	3
Embarked	418	3	S	270

```
[ ]: df.isna().sum()
```

```
[ ]: PassengerId      0
Survived             0
Pclass              0
Name                0
Sex                 0
Age                86
SibSp              0
Parch              0
Ticket             0
Fare               1
Cabin             327
Embarked           0
dtype: int64
```

```
[ ]: df['Cabin'].value_counts()
```

```
[ ]: B57 B59 B63 B66      3
B45                      2
C89                      2
C55 C57                  2
A34                      2
..
E52                      1
D30                      1
E31                      1
```

```
C62 C64          1
C105             1
Name: Cabin, Length: 76, dtype: int64
```

```
[ ]: #dropping irrelevant column
df.drop(['PassengerId', 'Name', 'Ticket'], axis=1, inplace=True)
df
```

```
[ ]:      Survived  Pclass    Sex   Age  SibSp  Parch    Fare  Cabin Embarked
0           0         3   male  34.5     0     0    7.8292   NaN      Q
1           1         3  female  47.0     1     0    7.0000   NaN      S
2           0         2   male  62.0     0     0    9.6875   NaN      Q
3           0         3   male  27.0     0     0    8.6625   NaN      S
4           1         3  female  22.0     1     1   12.2875   NaN      S
..          ...      ...    ...   ...    ...    ...    ...    ...
413          0         3   male   NaN     0     0    8.0500   NaN      S
414          1         1  female  39.0     0     0   108.9000  C105      C
415          0         3   male  38.5     0     0    7.2500   NaN      S
416          0         3   male   NaN     0     0    8.0500   NaN      S
417          0         3   male   NaN     1     1   22.3583   NaN      C
```

[418 rows x 9 columns]

```
[ ]: df['Age']=df['Age'].fillna(df['Age'].mean())
df['Fare']=df['Fare'].fillna(df['Fare'].mode()[0])
```

```
[ ]: df.isna().sum()
```

```
[ ]: Survived      0
Pclass           0
Sex              0
Age              0
SibSp           0
Parch           0
Fare             0
Cabin          327
Embarked         0
dtype: int64
```

```
[ ]: #drops the "Cabin" column from the DataFrame
df.drop(['Cabin'], axis=1, inplace=True)
```

```
[ ]: df.isna().sum()
```

```
[ ]: Survived      0
Pclass           0
Sex              0
```

```
Age          0
SibSp        0
Parch        0
Fare         0
Embarked     0
dtype: int64
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    418 non-null    int64
 1   Pclass      418 non-null    int64
 2   Sex         418 non-null    object
 3   Age         418 non-null    float64
 4   SibSp       418 non-null    int64
 5   Parch       418 non-null    int64
 6   Fare        418 non-null    float64
 7   Embarked    418 non-null    object
dtypes: float64(2), int64(4), object(2)
memory usage: 26.2+ KB
```

```
[ ]: df.dtypes
```

```
[ ]: Survived    int64
Pclass         int64
Sex            object
Age            float64
SibSp          int64
Parch          int64
Fare           float64
Embarked       object
dtype: object
```

## DATA VISUALIZATION


```
[ ]: df['Survived'].value_counts()
```

```
[ ]: 0    266
     1    152
     Name: Survived, dtype: int64
```

```
[ ]: plt.figure(figsize=(10,8))
     label=['not survived','survived']
```

```
plt.pie(df['Survived'].value_counts(),labels=label,autopct='%1.  
↪2f%%',explode=[0,0.1],shadow=True)  
plt.legend()
```


```
[ ]: <matplotlib.legend.Legend at 0x793198043010>
```



```
[ ]: df['Sex'].value_counts()
```

```
[ ]: male      266  
     female   152  
     Name: Sex, dtype: int64
```


```
[ ]: color=['teal','green']  
sns.countplot(x='Sex',data=df,hue='Sex')  
plt.show()
```



```
[ ]: sns.countplot(x='Pclass',data=df)
```


```
[ ]: <Axes: xlabel='Pclass', ylabel='count'>
```






```
[ ]: sns.barplot(x=df['Pclass'],y=df['Survived'])
```

```
[ ]: <Axes: xlabel='Pclass', ylabel='Survived'>
```



```
[ ]: sns.countplot(x='Pclass',data=df,hue='Survived')
```

```
[ ]: <Axes: xlabel='Pclass', ylabel='count'>
```



## CORRELATION

```
[ ]: df.corr()
```

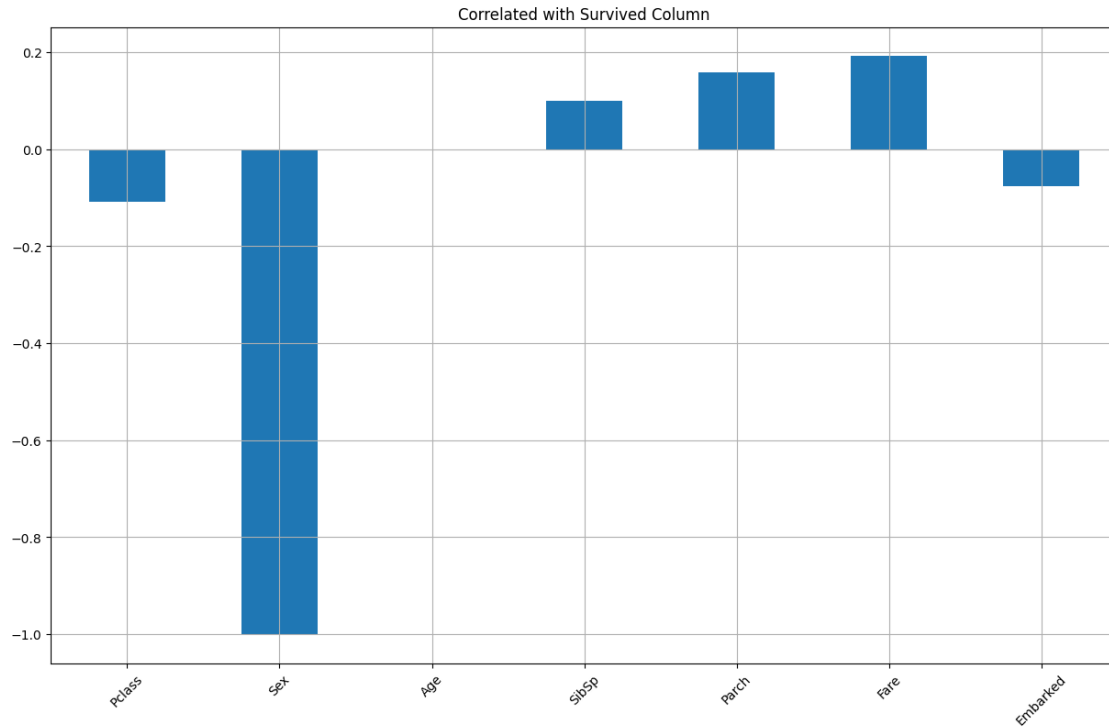
```
[ ]:
      Survived   Pclass    Age   SibSp   Parch   Fare
Survived  1.000000 -0.108615 -0.000011  0.099943  0.159120  0.192229
Pclass    -0.108615  1.000000 -0.440782  0.001087  0.018721 -0.577491
Age        -0.000011 -0.440782  1.000000 -0.079535 -0.045073  0.323839
SibSp      0.099943  0.001087 -0.079535  1.000000  0.306895  0.172034
Parch      0.159120  0.018721 -0.045073  0.306895  1.000000  0.230411
Fare       0.192229 -0.577491  0.323839  0.172034  0.230411  1.000000
```

```
[ ]: plt.figure(figsize=(14,8))
      sns.heatmap(df.corr(),annot=True)
      plt.show()
```



```
[ ]: df1=df.drop(columns='Survived')
df1.corrwith(df['Survived']).plot.bar(figsize=(15,9),title='Correlated with_
↳Survived Column',rot=45,grid=True)
```

```
[ ]: <Axes: title={'center': 'Correlated with Survived Column'}>
```



## LABEL ENCODING

```
[ ]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df['Sex']=le.fit_transform(df['Sex'])
df['Embarked']=le.fit_transform(df['Embarked'])
df
```

```
[ ]:
Survived  Pclass  Sex    Age    SibSp  Parch    Fare  Embarked
0         0      3     1  34.50000    0     0    7.8292      1
1         1      3     0  47.00000    1     0    7.0000      2
2         0      2     1  62.00000    0     0    9.6875      1
3         0      3     1  27.00000    0     0    8.6625      2
4         1      3     0  22.00000    1     1   12.2875      2
..      ...    ...    ...    ...    ...    ...    ...
413        0      3     1  30.27259    0     0    8.0500      2
414        1      1     0  39.00000    0     0   108.9000      0
415        0      3     1  38.50000    0     0    7.2500      2
416        0      3     1  30.27259    0     0    8.0500      2
417        0      3     1  30.27259    1     1   22.3583      0
```

[418 rows x 8 columns]

## X AND Y SEPARATION

```
[ ]: x=df.drop(columns='Survived',axis=1)
y=df['Survived']
```

```
[ ]: x
```

```
[ ]:      Pclass  Sex      Age  SibSp  Parch      Fare  Embarked
0         3    1  34.50000    0      0    7.8292         1
1         3    0  47.00000    1      0    7.0000         2
2         2    1  62.00000    0      0    9.6875         1
3         3    1  27.00000    0      0    8.6625         2
4         3    0  22.00000    1      1   12.2875         2
..      ...  ...      ...      ...      ...      ...
413        3    1  30.27259    0      0    8.0500         2
414        1    0  39.00000    0      0   108.9000         0
415        3    1  38.50000    0      0    7.2500         2
416        3    1  30.27259    0      0    8.0500         2
417        3    1  30.27259    1      1   22.3583         0
```

[418 rows x 7 columns]

```
[ ]: y
```

```
[ ]: 0    0
1    1
2    0
3    0
4    1
..
413   0
414   1
415   0
416   0
417   0
Name: Survived, Length: 418, dtype: int64
```

## SPLITTING IN TO TRAINING AND TESTING DATA

```
[ ]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪3,random_state=42)
```

## FEATURE SCALING

```
[ ]: from sklearn.preprocessing import StandardScaler
sd=StandardScaler()
x_train=sd.fit_transform(x_train)
x_test=sd.fit_transform(x_test)
```

```
[ ]: from sklearn.metrics import   
      accuracy_score, classification_report, ConfusionMatrixDisplay
```

## 2 MODEL BUILDING AND EVALUATION

K Nearest Neighbor Classifier

SVC

LOGISTIC REGRESSION

DECISION TREE CLASSIFIER

RANDOM FOREST CLASSIFIER

### 1.KNN (K Nearest Neighbours)

```
[ ]: #model creation   
from sklearn.neighbors import KNeighborsClassifier   
knn=KNeighborsClassifier(n_neighbors=7)   
knn.fit(x_train,y_train)   
knn_pred=knn.predict(x_test)
```


```
[ ]: #accuracy score   
knn_acc=accuracy_score(y_test,knn_pred)   
knn_acc
```

```
[ ]: 0.9841269841269841
```

```
[ ]: #classification report and confusion matrix   
print(classification_report(y_test,knn_pred))   
print(ConfusionMatrixDisplay.from_predictions(y_test,knn_pred))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	85
1	0.98	0.98	0.98	41
accuracy			0.98	126
macro avg	0.98	0.98	0.98	126
weighted avg	0.98	0.98	0.98	126

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at   
0x793198041150>
```



## 2.SVC (SUPPORT VECTOR CLASSIFIER)

```
[ ]: #model creation
from sklearn.svm import SVC
model=SVC()
model.fit(x_train,y_train)
y_pred_svm=model.predict(x_test)
```

```
[ ]: #accuracy score
svm_acc=(accuracy_score(y_test,y_pred_svm))
svm_acc
```

```
[ ]: 0.9920634920634921
```


```
[ ]: print(classification_report(y_test,y_pred_svm))
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred_svm))
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	85
1	1.00	0.98	0.99	41



accuracy			0.99	126
macro avg	0.99	0.99	0.99	126
weighted avg	0.99	0.99	0.99	126

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7931993cf670>



### 3.LOGISTIC REGRESSION

```
[ ]: #model creation
from sklearn.linear_model import LogisticRegression
lg=LogisticRegression(random_state=42)
lg.fit(x_train,y_train)
y_pred_lg=lg.predict(x_test)
```


```
[ ]: #accuracy score
lg_acc=(accuracy_score(y_test,y_pred_lg))
lg_acc
```

```
[ ]: 1.0
```

```
[ ]: #classification report and confusion matrix
print(classification_report(y_test,y_pred_lg))
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred_lg))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85
1	1.00	1.00	1.00	41
accuracy			1.00	126
macro avg	1.00	1.00	1.00	126
weighted avg	1.00	1.00	1.00	126

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x793197fc9090>



#### 4.DECISION TREE CLASSIFIER


```
[ ]: #model creation
from sklearn.tree import DecisionTreeClassifier
decision=DecisionTreeClassifier()
decision.fit(x_train,y_train)
```

```
y_pred_decision=decision.predict(x_test)
```

1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85
1	1.00	1.00	1.00	41
accuracy			1.00	126
macro avg	1.00	1.00	1.00	126
weighted avg	1.00	1.00	1.00	126

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x793197d94af0>




```
[ ]: #accuracy score
dtree_acc=accuracy_score(y_test,y_pred_decision)
dtree_acc
```

[ ]: 1.0

```
[ ]: #classification report and confusion matrix
print(classification_report(y_test,y_pred_decision))
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred_decision))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85
1	1.00	1.00	1.00	41
accuracy			1.00	126
macro avg	1.00	1.00	1.00	126
weighted avg	1.00	1.00	1.00	126

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7931980425c0>



## RANDOM FOREST CLASSIFIER

```
[ ]: #model creation
from sklearn.ensemble import RandomForestClassifier
random=RandomForestClassifier(random_state=42)
random.fit(x_train,y_train)
```

```
y_pred_random=random.predict(x_test)
```


```
[ ]: #accuracy score  
rf_acc=accuracy_score(y_test,y_pred_random)  
rf_acc
```

```
[ ]: 1.0
```

```
[ ]: #classification report and confusion matrix  
print(classification_report(y_test,y_pred_random))  
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred_random))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85
1	1.00	1.00	1.00	41
accuracy			1.00	126
macro avg	1.00	1.00	1.00	126
weighted avg	1.00	1.00	1.00	126

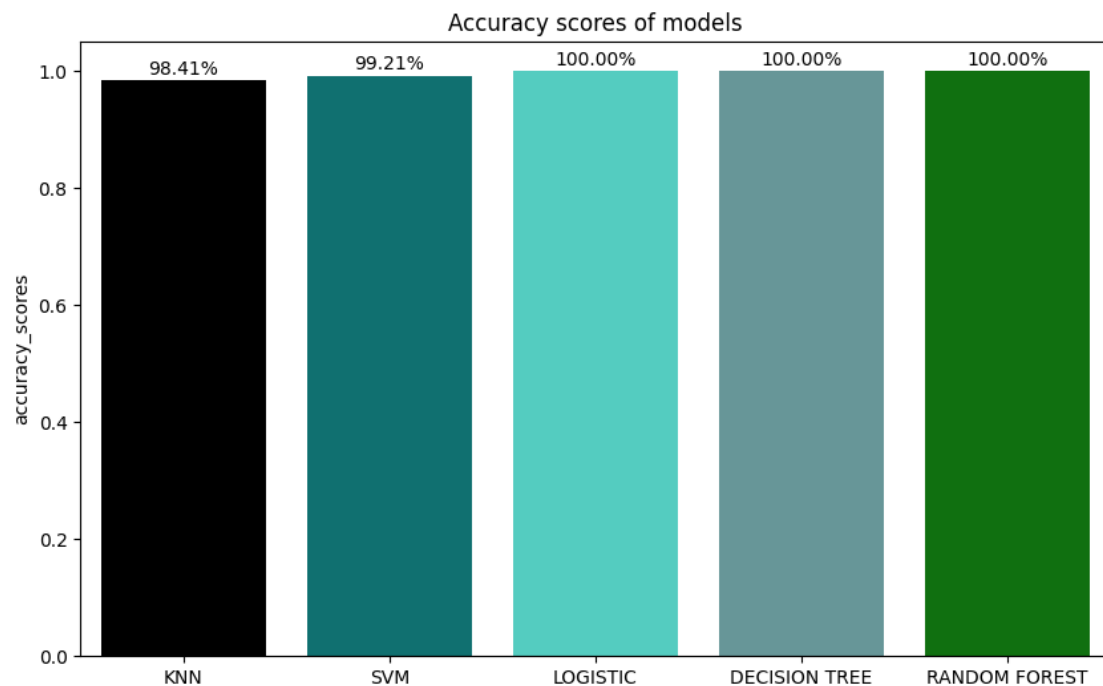
<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x793197762560>



```
[ ]: model=['KNN','SVM','LOGISTIC','DECISION TREE','RANDOM FOREST']
accuracy_scores=[knn_acc,svm_acc,lg_acc,dtree_acc,rf_acc]
accuracy_scores
```

```
[ ]: [0.9841269841269841, 0.9920634920634921, 1.0, 1.0, 1.0]
```

```
[ ]: color=['black','teal','turquoise','cadetblue','green']
plt.figure(figsize=(10,6))
sns.barplot(x=model,y=accuracy_scores,palette=color)
plt.ylabel('accuracy_scores')
plt.title('Accuracy scores of models')
# Adding percentage labels
for i, score in enumerate(accuracy_scores):
    plt.text(i, score +0.01, f'{score*100:.2f}%' , ha = 'center')
plt.show()
```



### 3 CONCLUSION

After a comprehensive analysis on the Titanic Survival Prediction dataset, multiple machine learning models were employed:

- **KNN:** 98.41% accuracy

- **SVM:** 99.21% accuracy
- **Logistic Regression:** 100% accuracy
- **Decision Tree:** 100% accuracy
- **Random Forest:** 100% accuracy

All models performed exceptionally well, achieving high accuracy scores. It's noteworthy that the Logistic Regression, Decision Tree, and Random Forest models achieved perfect accuracy, indicating robust predictive capabilities.